

NLP ISE 2026

PA1 - Text Classification

Due 25 August 2026

Instructions

- This is an individual assignment. You may discuss the tasks with classmates or the TA, but you must write your own code independently.
- This programming assignment consists of two parts:
 - Part 1 - You will implement the inference algorithm for logistic regression. Submit `logistic_regression.py` on MyCourseVille.
 - Part 2 - You will use sklearn library to train a logistic regression classifier and use an LLM library to implement a zero-shot LLM classifier.

Part 1 - Inference for Logistic Regression

We will implement the `TextClassifier` class, which loads a parameter matrix for a bag-of-words model. Here, every word found in the training set is used as a feature. For simplicity, this model does not include a bias (intercept) term in its calculations. You should first inspect the constructor and see what it does to the csv file.

The constructor will load the parameter matrix stored in a CSV file, where each row represents a word and its corresponding weights for different labels. The leftmost column contains words and is named "word". The remaining columns store weights, and their headers are the label names (these depend on the given CSV and are not always fixed as A, B, C—nor limited to just 3 labels).

Example:

word	A	B	C
money	0.2	-0.5	0.3
finance	0.4	0.3	0.2
dust	-0.2	-0.5	0.1
hate	0.1	0.1	0.1
mob	-0.3	0.6	-0.2

From this table, the weights for the word "dust" under labels A, B, and C are -0.2, -0.5, and 0.1, respectively.

Functions to Implement:

1. `get_all_possible_features`

Returns: A list of strings containing all features (words) in the model.

Example: `model.get_all_possible_features()` → ['money', 'finance', 'dust', 'hate', 'mob']

Hint: Each row's index in `self.model_parameters.index` corresponds to a word.

2. `get_all_possible_labels`

Returns: A list of strings of all possible labels (column headers except "word").

Example: `model.get_all_possible_labels()` → ['A', 'B', 'C']

3. `compute_probability`

Tokenize the text using NLTK's tokenizer. Use the parameter matrix to classify the text into one of the labels from the CSV. Calculate and return a dictionary of probabilities for each label, where:

Keys: Label names, Values: Corresponding probabilities

For example, the function should return something like: {'A': 0.32, 'B': 0.23, 'C': 0.43}

Hint: We can use this pseudocode to score the labels.

```
for label in labels:
    for feature in token_list:
        d.loc[feature][label] # weight for this feature-label pair
    update scores
```

4. `classify`

Predict the highest-probability label for the input text.

Example: For the probabilities {'A':0.32, 'B':0.23, 'C':0.43}, return 'C'

Part 2 - Model Comparison

The goal of this part is to compare the evaluation results from a logistic regression model and a zero-shot LLM classifier. The dataset (training and test sets) that we will use for this part is the financial news data.

You are required to answer Questions 1-5 and submit your work on mycourseville as pa-part2.pdf

2.1 Logistic Regression

Train a logistic regression model using bag-of-word features on the training set and evaluate the model on the test set. (You might have done this in class already).

Question 1: Report the precision, recall, F1-score for each label, as well as the overall accuracy. Interpret the model performance like we did in class.

Question 2: Identify the two best-performing labels. Examine the top 10 highest-weighted features for each of these labels and explain why the classifier performs well for them.

Question 3: Identify the two worst-performing labels. Examine the top 10 highest-weighted features for each of these labels and explain why the classifier does not perform well for them.

2.2 Zero-shot LLM Classifier

Implement a zero-shot LLM classifier for the same task and evaluate the performance.

1. **Prompt design:** Write a templated prompt (i.e. a prompt with a placeholder for the input text) that instructs an LLM to classify a news text string (or multiple at once) into one of the categories (labels) from the training set. The prompt should:
 - Define the role of the LLM in this task
 - Clearly explain each label
 - Specify the desired output format
 - Place the input text at the end of the prompt (this helps reduce token usage through context caching)

You should test your prompt on the web-based LLM such as ChatGPT, Gemini, or Claude to make sure that it works decently well.

Question 4: What is the prompt that you end up using? Include a screenshot from ChatGPT, Gemini, or Claude that you have tested the prompt a little bit.

2. **Implementation:** Write a function that takes a list of news text, inserts it into the prompt, calls an LLM API through a client library, and returns a label string.

The `typhoon-v2.5-30b-a3b-instruct` is recommended as it gives unlimited free api tokens (within rate limit). However, you may use other models if you prefer. You will need to create a free API key here: <https://playground.opentyphoon.ai/settings/api-key>. Of course, you are welcome to use OpenAI or Gemini models but that will cost you 1-2 USD.

Alternatively, you can also download and run inference on qwen models or other open-weight models that are no larger than 7 billion parameters on Colab. If a model is larger than that, Colab might not have enough memory to run it.

3. **Evaluation:** Randomly sample 200 instances from the test set, run the LLM with your prompt on them, collect the predicted labels, and evaluate them against the gold-standard labels. Evaluate the logistic regression that you trained in Part 2.1 on the same 200 instances.

Question 5: Which LLM do you use for this? Does your zero-shot LLM classifier work better or worse than the logistic regression? Give a possible explanation for this result.